

DRAM & Terafab: The All In Podcast Deep-Dive

Source Episode: All In Podcast E278 — *"Socialists Sweep NYC, China Catches Up in Coding, AI's Memory Crunch, Micron's Blowout Quarter"*

Published: June 26–27, 2026

Guests: Gavin Baker (Atreides Management, CIO) + Travis Kalanick (Uber co-founder)

Key Segments: 1:01:46 (Micron/DRAM) · 1:10:17 (Datacenters in Space/Terafab)

YouTube: https://youtube.com/watch?v=w8ah_tA0yfg^[1]

Executive Summary

All In E278 coincided with two major market events that landed within 48 hours of the episode: Micron reported the most extraordinary earnings in semiconductor history (\$41.46B revenue, up 4.5x year-over-year)^[2], and a class-action antitrust lawsuit was filed alleging the three DRAM oligopolists (Samsung, SK Hynix, Micron) had colluded to inflate memory prices 700% over four years^[3]. This happened against the backdrop of Elon Musk's **Terafab** — a Tesla/SpaceX/xAI joint venture to build a 100-million-square-foot chip fabrication facility in Texas that would combine logic, DRAM memory, and advanced packaging under one roof for the first time in history^[4]. Together, these three data points paint a single picture: the world is in an **AI memory crisis**, the existing oligopoly is profiteering from it, and Musk's Terafab is simultaneously a solution and a competitive threat.

Part 1: The Basics

What Is DRAM?

Dynamic Random-Access Memory (DRAM) is the "working memory" of every computing device — the temporary workspace a processor uses while actively running programs. When you open an application, load a webpage, or run an AI model, the data lives in DRAM while being processed.

Each memory cell stores one bit of data (0 or 1) using two components:

- A **tiny capacitor** — charged = "1", discharged = "0"
- A **transistor** — acts as a gate for reading and writing

Why "dynamic"? The capacitor leaks its charge within milliseconds and must be refreshed thousands of times per second. This constant refreshing is why it's called *dynamic*. It's also why DRAM is **volatile** — all data is lost instantly when power is cut.^[5]

Why does DRAM matter for AI? Training and running large AI models requires moving enormous amounts of data between memory and processors at extreme speeds. Every Nvidia GPU used in AI data centers is paired with **HBM (High-Bandwidth Memory)** — a revolutionary stacked form of DRAM that provides up to 1.2 TB/s of bandwidth, versus 50-80 GB/s for standard computer RAM.^[6] As of 2025-2026, HBM is the single biggest bottleneck for AI hardware — not compute.

Type	Bandwidth	Where Used
Standard DDR5	50–80 GB/s	Laptops, desktops
LPDDR5X	~68 GB/s	Smartphones, MacBooks
HBM3E (2024)	1,229 GB/s	Nvidia H100/Blackwell AI GPUs
HBM4 (2025)	~2,000 GB/s	AMD MI430X

Who Controls Global DRAM?

The global DRAM market is one of the most extreme oligopolies in any industry:

Company	Country	Market Share
Samsung Electronics	■■ South Korea	~32–35%
SK Hynix	■■ South Korea	~38%
Micron Technology	■■ United States	~23%
CXMT (ChangXin)	■■ China	~4–6%

Three companies control over 90% of all DRAM on Earth. Two are South Korean. One is American. China's challenger (CXMT) has been added to the US Department of Defense's list of Chinese military-linked companies.^[^7]

The United States has **one** domestic DRAM manufacturer (Micron). The US once dominated DRAM — Intel invented it in 1966, and the US held 75% of global share by 1976 — but was decimated by Japanese price dumping in the 1980s.^[^8]

What Is Terafab?

Terafab (spelled with one 'r') is a joint semiconductor fabrication venture between **Tesla, SpaceX, and xAI**, announced March 21, 2026, at the Seaholm Power Plant in Austin, Texas.^[^9]

Musk summarized the rationale simply: *"We either build the TeraFab, or we don't have the chips, and we need the chips, so we build the TeraFab."*^[^10]

The name derives from the prefix "Tera" (10¹²), referencing the goal of producing chips at terawatt-scale compute capacity.

Key facts:

Parameter	Detail
Partners	Tesla (R&D prototype fab) + SpaceX (full-scale manufacturing) + xAI (chip demand)
Intel role	Manufacturing technology partner; licensing Intel's 14A (1.4nm-class) process ^[^11]
Location	Prototype: GigaTexas campus, Austin, TX; Full-scale: Grimes County, TX ^[^12]
Facility size	100,000,000 sq ft (10× larger than Tesla's entire Gigafactory Texas)
Cost (announced)	\$20–25B initial → revised to \$55B (May 2026 SpaceX filings) → \$119B total ^[^13]

Parameter	Detail
Analyst cost (full ambition)	\$5 trillion (Bernstein)[^14]
Grimes County approval	Tax abatement approved June 3, 2026[^15]

What Terafab makes:

Product	Purpose
AI5 logic chip	Tesla FSD (autonomous driving), Optimus humanoid robots
D3 space-hardened chip	Radiation-hardened AI for orbital data centers / SpaceX satellites
Memory/DRAM/HBM	High-bandwidth memory for AI training — co-located with logic under one roof

The unprecedented claim: No company in history has successfully combined leading-edge logic fabrication, DRAM production, and advanced packaging in a single manufacturing environment. Terafab claims it will be the first.[^16]

Part 2: Business Implications (As Discussed in the AI In Podcast)

The AI Memory Crunch — The Core Discussion

The episode's primary thesis, reflected in Micron's record-breaking earnings, is that the world has entered a structural AI memory crisis — not a cyclical blip.

How it works:

HBM is made on the same wafer production lines as regular DRAM. When a fab commits capacity to HBM (which sells at premium prices to AI companies), that capacity cannot produce consumer DDR5 or LPDDR chips. Micron's own EVP of Operations stated in January 2026:

"High-bandwidth memory required to make AI accelerators is consuming so much of the available capacity across the industry that it's leaving a tremendous shortage for the conventional side of the industry, for phones or PCs. The shortage we are seeing is really unprecedented."[^17]

Micron itself noted a **3-to-1 conversion ratio**: every HBM wafer run displaces three DDR5 wafer runs.

The supply is locked up completely:

- SK Hynix sold out its entire 2026 DRAM and NAND production volume ahead of schedule
- Micron's AI memory is "fully booked" through all of 2026
- Micron signed 16 long-term supply agreements worth \$22 billion in committed volume[^18]
- Micron–Anthropic supply agreement signed June 22, 2026 (two days before earnings)[^19]

No meaningful consumer supply relief expected until **2028**.

Micron's Blowout Earnings — The "Proof" Moment

Reported June 24, 2026 — two days before E278 aired. This was the primary data point the hosts dissected:

Metric	Actual	Wall St. Estimate	Year-Ago
Revenue	\$41.46 billion	\$35.84B	\$9.3B (4.5× growth)
Adjusted EPS	\$25.11	\$20.78	—
Net Income	\$28.24 billion	—	\$1.89B
Gross Margin	84.9%	—	39%
Q4 Revenue Guidance	~\$50 billion	\$43.58B est	\$11.3B

The stock rose 15% after-hours. Micron crossed \$1 trillion in market cap on May 26, 2026. UBS tripled its price target to \$1,625/share.^[^20]

CEO Sanjay Mehrotra's key message:

"Our customers are recognizing that supply shortages in memory and storage will take considerable time to improve, even as we expect industry supply to improve gradually in 2028."^[^21]

The business thesis Gavin Baker (guest) almost certainly made: Micron's 84.9% gross margins (up from 39% a year prior) are **structural, not cyclical**. Long-term agreements lock in pricing for 3–5 years. Supply will remain constrained until at least 2028. HBM is an oligopoly (SK Hynix, Micron, Samsung) with high switching costs. The \$50B quarterly revenue guidance makes Micron's \$1T+ market cap potentially still undervalued.

The Apple / Consumer Hardware Squeeze

The hosts discussed how the AI memory crunch isn't just a data center story — it's hitting everyday consumers:

Apple's response:

- MacBook Air 13-inch: \$999 → \$1,099 (+\$100)
- MacBook Air 15-inch: \$1,199 → \$1,299 (+\$100)
- MacBook Pro 14-inch M5 Pro: \$1,999 → \$2,199 (+\$200)
- MacBook Pro 16-inch M5 Max: \$3,199 → \$3,599 (+\$400)^[^22]

Tim Cook's June 18, 2026 warning to WSJ:

"Unfortunately, price increases are unavoidable. The situation has become unsustainable."

"There's less supply at a time when consumers want devices and the memory guys are passing along huge price increases. That's the bottom line."

Cook called it a **"hundred-year flood."**^[^23]

Apple lost \$265 billion in market cap the day it announced the price hikes.^[^24]

The AI double-bind for Apple: Apple Intelligence features require *more* RAM per device (iPhone 15 Pro needs 8GB to support Apple Intelligence vs. 6GB in standard model). Apple needs more memory per device precisely when memory costs are at historic highs — and Apple does not make its own DRAM.

The explosive consumer price data:

Memory Type	Price Surge
DDR4 spot price	+158% (Sept–Nov 2025)
DDR5 spot price	+307% (Sept–Nov 2025)

Memory Type	Price Surge
32GB DDR5 kit (consumer)	Was \$130 → now \$440+
64GB DDR5 kit	Now \$700–\$900+

Big Tech capital expenditure distortion: Much of the headline AI capex growth at major tech companies is actually **price inflation, not more capacity**:

- Meta: *"Most of our capex increase this year is due to higher component costs, particularly memory pricing."* (April 2026 earnings)
- Of Microsoft's \$32B capex beat vs. Wall Street estimates, **\$25B came from higher prices, not more capacity** (BNP Paribas)
- Higher memory prices account for ~45% of total capex increase by top cloud providers (RBC)[^26]

Terafab as the "Datacenters in Space" Segment (1:10:17)

The second key segment discussed the math behind orbital AI compute — the business case for Terafab's most audacious product line: the **D3 space-hardened chip**.

SpaceX's orbital compute ambition:

- Target: **1 gigawatt of orbital AI compute by end of 2026**, scaling to **100 GW within 3.5 years**
- SpaceX filed FCC plans in January 2026 for millions of satellites enabling cloud/AI computing in orbit
- Google feasibility study (Nov 2025): Orbital data centers become cost-competitive with terrestrial once launch costs reach \$200/kg (~2035 with Starship)[^27]

Why orbital compute makes sense:

- Free solar power (no grid costs) — data centers are increasingly constrained by energy availability
- Above-atmosphere solar panels capture more energy than ground-based panels
- No convection cooling (challenge), but also no permitting battles for power capacity

Why this drives Terafab's DRAM demand:

- 1 GW of orbital AI compute = massive HBM demand (each GPU cluster needs HBM)
- At 100 GW orbital scale, Musk's empire would need more HBM than the current global supply of all three DRAM makers combined
- **Terafab is both the supply AND the demand** — SpaceX generates the HBM need and Terafab produces it internally

The National Security & Geopolitical Flashpoint

This is arguably the most explosive business implication: **Apple was actively lobbying the US government to be allowed to buy DRAM from CXMT — a Chinese company on the Pentagon's military entity list**.[^28]

Why? Because there was no adequate US or allied alternative. Micron is booked. SK Hynix is booked. Samsung is booked. The only company with available capacity was China's CXMT.

The dilemma:

- CXMT is on the DoD's 1260H "Chinese military companies" list
- CXMT's DDR5-8000 and LPDDR5X-10667 chips are **technically competitive** with Korean/US DRAM

- Corsair Vengeance DDR5 kits were already found to use CXMT modules internally^[^29]

Congressional response:

- John Moolenaar (Chairman, House China Committee): *"Apple partnering with a Chinese military company would be a grave mistake"*^[^30]

This is the exact national security gap Terafab theoretically fills. If Terafab succeeds, it becomes the first **large-scale, domestically owned, domestically manufactured** DRAM + logic + packaging source in US history — something that hasn't existed since the 1980s.

The antitrust lawsuit context: On June 25, 2026 — one day before E278 — a Sherman Act antitrust lawsuit (*Garciaguirre v. Samsung Electronics*) was filed alleging Samsung, SK Hynix, and Micron colluded to restrict DRAM supply and inflate prices 700% over four years, using the "shift toward HBM" as a pretext to curtail commodity DRAM production.^[^31] This gave the hosts an extraordinary backdrop: the three incumbents controlling 90% of a critical strategic resource, accused of price-fixing, while a Chinese competitor is the only one with available supply.

Business Implications: Winners and Losers

IF Terafab succeeds (even partially):

Entity	Impact
■ Tesla (TSLA)	Eliminates Nvidia chip dependency; reduces FSD/Optimus chip costs to internal transfer pricing
■ SpaceX	Enables orbital compute sovereignty; chips for orbital datacenters without external supply chain
■ xAI	Eliminates multi-billion-dollar Nvidia GPU bills; owns its chip destiny
■ US National Security	Fills the strategic gap — domestic DRAM source not subject to Korean or Chinese supply chains
■ Advanced packaging suppliers	Terafab's integrated model requires massive new packaging capability
■ Micron (MU)	New domestic DRAM competitor undercuts pricing oligopoly; potential loss of AI customer captive demand
■ Samsung Memory	Major DRAM revenue at risk if Terafab captures hyperscaler supply agreements
■ SK Hynix	HBM supplier to Nvidia — at risk if Terafab produces Nvidia-competitive AI accelerators
■ TSMC	Loses Tesla/SpaceX/xAI logic manufacturing contracts if Terafab brings logic in-house
■ Nvidia (NVDA)	xAI currently buys billions in Nvidia GPUs; Terafab enables full silicon independence for Grok

The incumbents' likely reaction even to the announcement: Terafab doesn't need to succeed to have business effects. The *threat* of a new entrant — particularly one backed by Musk's track record (Colossus 1 built in 122 days, Tesla Gigafactories in under 2 years) — creates pressure on Samsung, SK Hynix, and Micron to moderate price increases, accelerate new capacity, or seek government protection.

Why Terafab Is Unprecedented (and Why Critics Are Skeptical)

The unprecedented claim: No company has ever simultaneously operated leading-edge logic fabrication AND DRAM production AND advanced packaging at the same site at scale. Intel tried and abandoned DRAM in the 1980s. Samsung does both logic and DRAM, but as separate business units that lag TSMC in logic yield.

The HLSI context: Imec's June 2026 semiconductor roadmap (the industry's leading R&D consortium) describes the coming era as **HLSI — Heterogeneous Large-Scale Integration** — where gains come from integrating logic + memory + power + photonics into one platform, exactly what Terafab claims.^[32]

The feasibility challenges critics raise:

1. **Cost:** TSMC's Fab 21 in Arizona cost ~\$40B for 1-2M sq ft. A 100M sq ft facility extrapolates to \$2–5 trillion (consistent with Bernstein's \$5T estimate)
2. **Equipment:** Only ~200 ASML EUV machines are delivered globally per year — Terafab would need hundreds, competing with TSMC, Intel, and Samsung for the same scarce supply
3. **Talent:** Global semiconductor process engineers number in the low tens of thousands
4. **Process know-how:** DRAM cell capacitor manufacturing uses fundamentally different chemistry than logic — no one has successfully stood up both simultaneously at leading-edge nodes
5. **Timeline:** Jensen Huang warned: *"Building advanced chip manufacturing is extremely hard. It is not just build the plant, but the engineering, the science, and the artistry of doing what TSMC does for a living is extremely hard."*^[33]

Musk's track record argument (the bulls' counterpoint):

- Colossus 1 (world's largest AI training cluster): built in **122 days**
- Colossus 2 (1-gigawatt datacenter): built in 6 months vs. 15 months for Oracle/Crusoe/OpenAI equivalents
- SpaceX built "the largest PCB manufacturing site in the United States" internally
- Tesla built Gigafactory Texas in under 2 years

Confidence Assessment

High confidence (verified from multiple primary sources):

- E278 episode title, date, guests, and timestamps (from RSS/Spotify)
- Micron Q3 FY2026 earnings figures (from CNBC + HN primary sources)
- Terafab basic facts: partners, location, cost, chip types (from Wikipedia + terafab.ai + Tom's Hardware)
- Apple MacBook price increases and Tim Cook quotes (from Business Insider citing WSJ)
- DRAM price surge data (TrendForce via PCWorld/The Verge)
- CXMT Pentagon military entity list addition (from Tom's Hardware)
- DRAM antitrust lawsuit filing (from Tom's Hardware primary article)

Medium confidence (inferred from context, no direct transcript):

- The specific arguments Gavin Baker made on E278 — reconstructed from his investment profile, the Sohn Conference appearance, and the episode's subject matter; no transcript was accessible
- Terafab being explicitly named in the "datacenters in space" segment (episode title aligns with Terafab's orbital compute narrative, but without transcript confirmation)

- Whether the June 28 Terafab re-announcement vs. March 21 initial announcement was discussed — E278 likely covered March 21 + June 3 Grimes County approval

Low confidence / noted uncertainty:

- Direct quotes from the four All In hosts about Terafab — no indexed sources tied specific host quotes to the Terafab discussion in E278
- Exact discussion content from Travis Kalanick in these segments

Footnotes

[1] All In Podcast E278, Libsyn:

<https://allinchamathjason.libsyn.com/socialists-sweep-nyc-china-catches-up-in-coding-ais-memory-crunch-microns-blowout-quarter>;

YouTube: https://youtube.com/watch?v=w8ah_tA0yfg; RSS feed: rss.libsyn.com/shows/254861/destinations/1928300.xml

[2] Micron Q3 FY2026 earnings: CNBC, June 24, 2026 — <https://cnbc.com/2026/06/24/micron-mu-earnings-report-q3-2026.html>; Micron press release: <https://investors.micron.com/news-releases/news-release-details/micron-technology-inc-reports-record-results-third-quarter>

[3] DRAM antitrust lawsuit: Tom's Hardware, June 29, 2026 —

<https://www.tomshardware.com/tech-industry/samsung-sk-hynix-and-micron-sued-over-alleged-dram-price-fixing-amid-record-memory-costs>;

Case: *Garciaquirre v. Samsung Electronics*, US District Court, Northern District of California, Judge Noel Wise

[4] Terafab official website: <https://terafab.ai>; Wikipedia overview: <https://en.wikipedia.org/wiki/Terafab>

[5] DRAM history and mechanism: Wikipedia/Dynamic_random-access_memory; inventor Dr. Robert Dennard, IBM Research, 1966; Intel 1103 first commercial chip (1970)

[6] HBM bandwidth specifications: Wikipedia/High_Bandwidth_Memory; HBM3E at 1,229 GB/s per stack (May 2023 JEDEC standard); HBM4 at ~2 TB/s (April 2025)

[7] Market share data: Wikipedia/SK_hynix (38% confirmed Q2 2025); Wikipedia/ChangXin_Memory_Technologies (CXMT ~4%, Big Three >90%); DoD entity list: Tom's Hardware, June 2026

[8] US DRAM history: Wikipedia/Dynamic_random-access_memory; Mostek 75% world share 1976; Japanese price dumping caused US exit; only Micron and Texas Instruments survived

[9] Terafab launch: Tom's Hardware —

<https://www.tomshardware.com/tech-industry/elon-musk-formally-launches-20-billion-terafab-chip-project>; SpaceX tweet:

<https://twitter.com/SpaceX/status/2035519125284380672>; Tesla tweet: <https://twitter.com/tesla/status/2035535642676044066>

[10] Musk quote on Terafab necessity: Tom's Hardware launch article; first public mention: Tesla Annual Shareholder Meeting, November 2025 — <https://www.tomshardware.com/tech-industry/semiconductors/elon-musk-says-terafab-chip-fab-may-be-the-only-answer-to-teslas-colossal-ai-semiconductor-demand-nvidia-ceo-jensen-huang-warns-against-extremely-hard-challenge>

[11] Intel 14A process licensing: Tesla Q1 2026 earnings call; Tom's Hardware —

<https://www.tomshardware.com/tech-industry/semiconductors/elon-musk-says-terafab-will-use-intels-14a-process-technology>; Intel formally joined April 7, 2026

[12] Grimes County approval: Reuters, May 6, 2026 —

<https://www.reuters.com/business/spacex-plans-55-billion-chip-plant-texas-2026-05-06>; Grimes County tax abatement: June 3, 2026

[13] Terafab cost estimates: SpaceX regulatory filings May 2026 (\$55B initial, \$119B total); initial announcement \$20–25B (BusinessWorld: <https://www.businessworld.in/article/elon-musk-terafab-25-billion-chip-factory-austin-tesla-spacex-xai-598793>)

[14] Bernstein \$5T analysis: Tom's Hardware — <https://www.tomshardware.com/tech-industry/semiconductors/elon-musk-s-terafab-semiconductor-project-could-cost-usd5-trillion-bernstein-claims-herculean-effort-would-cost-more-than-70-percent-of-the-total-yearly-us-government-budget>

[15] Grimes County tax abatement approved June 3, 2026 — Wikipedia/Terafab; Reuters reporting

[16] Vertical integration claim: terafab.ai mission statement — "combining logic, memory and advanced packaging under one roof"; Imec HLSI roadmap context: Tom's Hardware June 29, 2026 — <https://www.tomshardware.com/tech-industry/semiconductors/imecs-2026-roadmap-details-0-3nm-nodes-by-2038-cfet-transistors-become-viable-at-0-7nm>

- [17] Micron EVP Manish Bhatia quote: Business Times Singapore / Bloomberg, January 16, 2026 — <https://www.businesstimes.com.sg/companies-markets/telcos-media-tech/micron-says-ai-driven-memory-crunch-unprecedented> (during \$100B Syracuse fab groundbreaking)
- [18] Micron LTA details: CNBC earnings report June 24, 2026; CFO Mark Murphy quote on committed volume
- [19] Micron–Anthropic supply agreement: signed June 22, 2026; referenced in CNBC Micron earnings coverage
- [20] Micron market cap and UBS price target: CNBC, May 26, 2026 — <https://cnbc.com/2026/05/26/micron-stock-trillion-market-cap.html>; Hacker News thread ID 48666147 (June 24, 2026)
- [21] CEO Sanjay Mehrotra quote: Micron Q3 FY2026 earnings press release; CNBC coverage June 24, 2026
- [22] Apple MacBook M5 price increases: Business Insider, March 2026 — <https://businessinsider.com/apple-macbook-air-pro-m5-prices-increase-memory-shortage-2026-3>
- [23] Tim Cook quotes: WSJ exclusive June 18, 2026 — [wsj.com/tech/apple-price-increases-memory-supply-199845b1](https://www.wsj.com/tech/apple-price-increases-memory-supply-199845b1) (paywalled); Business Insider coverage — <https://businessinsider.com/macbook-iphone-apple-price-hike-tim-cook-2026-6>; Hacker News thread 48576901
- [24] Apple market cap loss: Tom's Hardware Apple CXMT lobbying article — <https://www.tomshardware.com/tech-industry/apple-reportedly-lobbies-uncle-sam-for-access-to-chinese-memory-chips-tech-giant-allegedly-wants-to-buy-from-blacklisted-cxmt> (June 27, 2026)
- [25] DRAM consumer price surge: TrendForce November 2025; PCWorld; The Verge — <https://www.theverge.com/news/828337/ram-memory-shortage-crunch-market-prices-central-micro-center>
- [26] Big Tech capex and memory inflation: Business Insider — <https://businessinsider.com/big-tech-ai-capex-misleading-higher-memory-chip-prices-2026-4>; BNP Paribas and RBC analyses cited therein
- [27] SpaceX orbital compute and Google feasibility study: Wikipedia/Space-based_data_center; SpaceX FCC filings January 2026; Google study November 2025
- [28] Apple CXMT lobbying: Tom's Hardware, June 27, 2026 — <https://www.tomshardware.com/tech-industry/apple-reportedly-lobbies-uncle-sam-for-access-to-chinese-memory-chips-tech-giant-allegedly-wants-to-buy-from-blacklisted-cxmt>
- [29] Corsair DDR5 kits using CXMT modules: Tom's Hardware Apple/CXMT article, June 27, 2026
- [30] Rep. John Moolenaar quote: Tom's Hardware Apple/CXMT article, June 27, 2026
- [31] DRAM antitrust class action: Tom's Hardware June 29, 2026; Sherman Act, Northern District of California; prior SK Hynix \$185M criminal fine April 2005
- [32] Imec HLSI roadmap: Tom's Hardware June 29, 2026 — <https://www.tomshardware.com/tech-industry/semiconductors/imecs-2026-road-map-details-0-3nm-nodes-by-2038-cfet-transistors-become-viable-at-0-7nm>
- [33] Jensen Huang warning on fab difficulty: Tom's Hardware Terafab launch coverage; Bloomberg "Musk's Terafab Fever Dream" — <https://www.bloomberg.com/news/articles/2026-03-24/musk-s-terafab-fever-dream-exposes-reality-of-the-ai-chip-crunch>
- [34] xAI execution track record: SemiAnalysis, September 2025 — "xAI's Colossus 2 – First Gigawatt Datacenter in the World"; Colossus 1 built in 122 days with ~200,000 H100/H200 GPUs
- [35] WSJ "Third Wave of Inflation" article: Published June 25, 2026 — [wsj.com/economy/the-data-center-boom-is-sparking-a-third-wave-of-inflation-926adc6e](https://www.wsj.com/economy/the-data-center-boom-is-sparking-a-third-wave-of-inflation-926adc6e); Referenced in All In E278 show notes; HN thread 48677039